

# Enterprise Execution Readiness

An Executive Brief · July 2026

*When software can act on its own, can your systems stop a bad action before it happens? Eleven enterprise platforms, compared in plain language.*

**About this brief.** Prepared by Mimir Labs Research. Mimir Labs is one of the eleven platforms assessed here, and readers should weigh that context. Scores are directional analyst judgments based on public product documentation as of July 31, 2026 — not vendor-certified tests. The full 32-page report, methodology, and complete 30-condition score matrix are available on request. All trademarks belong to their owners.

## 1. The question every executive will soon be asked

For fifty years, business software mostly kept score. It recorded orders, balances, inventory, and approvals — and people decided what should happen next. When the software was wrong, or a rule was missing, a person usually caught the problem before it became real.

That era is ending. AI assistants and automated connections between systems no longer just recommend. They act: they approve, release, post, ship, and pay — at machine speed, around the clock, across every system they can reach.

That changes the question boards should ask. Not “*Is our AI smart?*” or “*Is our data clean?*” but:

**When anything — a person, a program, or an AI — tries to make a change it shouldn’t, does our system refuse before the change happens? Or do we find out afterward?**

We call the ability to answer “it refuses — every time, on every path” **execution readiness**. This brief summarizes how eleven enterprise platforms compare on it, and gives you the questions to put to any vendor — including us.

**The one rule to remember.** A watchdog standing next to your systems can warn you. Only a checkpoint that every change must pass through can protect you. If even one side door exists — an administrator's tool, a data import, a connected system — the guarantee does not exist.

## 2. Why this matters now

The risk is not malicious AI. A well-built AI agent is an eager optimizer: give it a goal and access, and it will find the fastest available path — including paths your control designers never thought of as paths. No bad intent is required; obedience to the goal is enough. The common failure patterns are easy to picture:

- **Acting on stale information.** An assistant approves a payment against a budget that changed seconds earlier.
- **Slicing.** A forbidden \$250,000 commitment becomes five permitted \$50,000 steps.
- **Side doors.** The order screen refuses, so the change arrives through a data import or integration that never checks the rule.
- **Racing.** Two systems update the same order at the same moment — and both “succeed.”
- **Finding out tomorrow.** Monitoring flags the invalid change on the next report — after the money moved or the truck left.

Human organizations absorbed these weaknesses for decades because people worked at human speed and used judgment. Autonomous software does neither. The more your AI is allowed to do, the more the architecture of your systems — not the intelligence of the AI — determines your risk.

## 3. What we measured

The Enterprise Execution Readiness Index 2026 assessed eleven platforms — SAP, Oracle, Microsoft, Salesforce, Infor, Epicor, Sage X3, Acumatica, Odoo, Rappit, and Mimir Labs — across thirty conditions, scored 0–20. The thirty conditions boil down to five plain questions:

- **Do all your systems mean the same thing?** One shared business vocabulary — a single definition of “customer,” “order,” or “complete” that systems are required to use, not just encouraged to document.
- **Can data, changes, and rules move safely?** Migrations that preserve meaning, synchronization that respects your rules, and whether those rules are a portable company asset or trapped inside one product.

- **Can it run your business?** Depth of workflows, transactions, and manufacturing; identity and access; resilience — and whether customizations can quietly create side doors.
- **How strong are the AI tools — and the human controls?** The quality of assistants and agents, the business context they see, whether people can intervene, and whether you can reconstruct what an AI actually did.
- **Can anything slip past the final checkpoint?** The decisive question: is there one checkpoint where the rules are checked against current facts before any change becomes official — with a permanent record, and no exceptions for anyone?

## 4. What we found

**Every major platform is genuinely good — at different things.** SAP and Oracle run the deepest business operations in the world. Microsoft and Salesforce have the best AI tools and the broadest oversight of AI activity. Infor, Epicor, Sage X3, Acumatica, and Odoo run real factories and mid-market businesses every day. Rappit builds custom systems using AI developers — promising, and less proven. Mimir Labs is the only platform in the group designed around the final checkpoint itself.

**Being big is not the same as being closed.** Deep operations do not imply a single checkpoint. In every established platform, the protections are real but spread across many tools, screens, and configurations — which means their strength depends on how completely each customer assembled them, and side doors are possible.

**Vendors' AI safety features govern the AI — not the outcome.** Today's AI controls decide what an AI may see, say, and access. Far fewer decide what any actor — AI, human, or program — is allowed to make official. Both matter. Only the second is a guarantee.

**Averages hide the story; the shape is the finding.** Overall averages cluster in a narrow band (14.0–18.0 out of 20) and conceal opposite architectures. The domain profile is what to read:

| Platform          | Shared meaning | Safe movement | Runs the business | AI & oversight | Final checkpoint |
|-------------------|----------------|---------------|-------------------|----------------|------------------|
| <b>Mimir Labs</b> | <b>19.0</b>    | 17.2          | 18.0              | 16.0           | <b>20.0</b>      |
| Oracle            | 17.5           | 16.8          | <b>19.0</b>       | 19.0           | 17.2             |
| SAP               | 17.8           | 17.2          | 18.7              | 18.7           | 15.7             |
| Microsoft         | 17.5           | <b>18.0</b>   | 17.0              | <b>19.5</b>    | 13.8             |
| Salesforce        | 17.7           | 17.3          | 16.0              | <b>19.5</b>    | 14.5             |
| Infor             | 16.0           | 16.2          | 18.0              | 17.0           | 15.0             |
| Epicor            | 14.3           | 14.0          | 17.7              | 16.2           | 15.7             |
| Acumatica         | 13.5           | 13.8          | 16.7              | 16.0           | 14.8             |
| Rappit            | 12.5           | 14.2          | 15.8              | 17.3           | 14.7             |
| Sage X3           | 13.5           | 13.5          | 16.7              | 15.5           | 14.2             |
| Odoo              | 12.0           | 13.2          | 15.8              | 15.7           | 13.5             |

Domain scores, 0–20, from the full 30-condition assessment. Bold marks the leader in each column. “Final checkpoint” aggregates the six conditions that determine whether an invalid change can become official.

**Read the last column with care.** It is the one place the market cleanly divides: every established platform lands between 13.5 and 17.2 — strong controls, assembled from parts — while Mimir Labs is at 20.0 because the checkpoint is the product. In fairness, the reverse is also visible: Mimir Labs’ lowest marks (connector count, AI-agent tooling) are the incumbents’ highest, and ours are deliberate — we do not ship connections that bypass the checkpoint, and we govern AI agents rather than build them.

## 5. Ten questions to ask any vendor — including us

Bring these to any demo or renewal conversation. The answers separate protection that is built in from protection that is configured, hoped for, or delegated to an implementation partner.

1. Where, exactly, is the official record of our business created?
2. Can anyone — an administrator, a data import, a connected system — change that record around your controls?
3. Are the rules checked before a change is saved, or discovered in reports afterward?

4. Is the check made against current facts at the moment of the change — not against a snapshot?
5. Could several individually allowed steps add up to a forbidden outcome?
6. What happens when two people — or two systems — change the same thing at the same time?
7. Are illegal status changes impossible, discouraged, or merely detected later?
8. Is the history complete and tamper-proof across every path — including bulk tools and administrators?
9. Can you show precisely why an action was allowed — the rule, its version, and the evidence behind it?
10. Which protections are built in, which are configured, and which depend on the implementation partner?

## 6. Using this in a buying decision

- **If your AI only reads, drafts, and recommends** — your existing governance, approvals, and access controls may well be sufficient. Move deliberately, not fearfully.
- **If your AI (or any automation) will approve payments, release inventory, change production, close quality events, or alter master data** — treat the final-checkpoint conditions as pass/fail requirements, not preferences. A brilliant score everywhere else cannot compensate for a side door.
- **Weight the domains to your situation.** A regulated manufacturer should weight operations and the checkpoint most heavily; an AI-program leader, the AI domain plus the checkpoint; a company untangling years of system sprawl, shared meaning and safe movement.

## 7. About Mimir Labs, and about this research

Mimir Labs builds Yggdrasil, an enterprise platform designed around the final checkpoint: every change — from a person, a program, or an AI — passes one gate that checks what it means, who authorized it, what evidence supports it, and whether it is legal given the current facts, before it becomes official. Every decision is permanently recorded and explainable. By design, we are not an AI-agent development platform and we do not offer a thousand-connector catalog: every connection we ship is governed, because one ungoverned connection would break the promise that justifies everything else.

**Research integrity.** We wrote this research and we appear in it; weigh that as you should. Three commitments back it: we credit competitors' real strengths by name and score ourselves lowest where the trade is real; we publish, for every vendor including ourselves, the evidence

that would change the scores; and we are building a public test suite of the failure patterns in Section 2 — unauthorized changes, sliced transactions, stale-information actions, side-door writes — which we will run on ourselves first, and publish.

Learn more: [mimirlabs.net](https://mimirlabs.net) · Full report and 30-condition matrix available on request.